

ACCELERATOR BASED PROGRAMMING UPPSALA UNIVERSITY FALL

ASSIGNMENT 2: Programming with CUDA

I have performed runs on the HSE university HPC-cluster (<https://hpc.hse.ru/en/hardware/hpc-cluster/>). The nodes are NVIDIA Tesla v100 (<https://www.nvidia.com/en-us/data-center/v100/>) nodes. Memory bandwidth is 900 GB/s and the performance is 15.7 teraFLOPS for Single-Precision.

Our general goal is to compute matrix-matrix product. In case of two matrices with $M \times N$ and $N \times K$ elements the result matrix has $M \times K$ elements.

To measure memory bandwidth for matrix-vector product (when $K=1$), we use $(M \times N + N \times K + M \times K)$ factor.

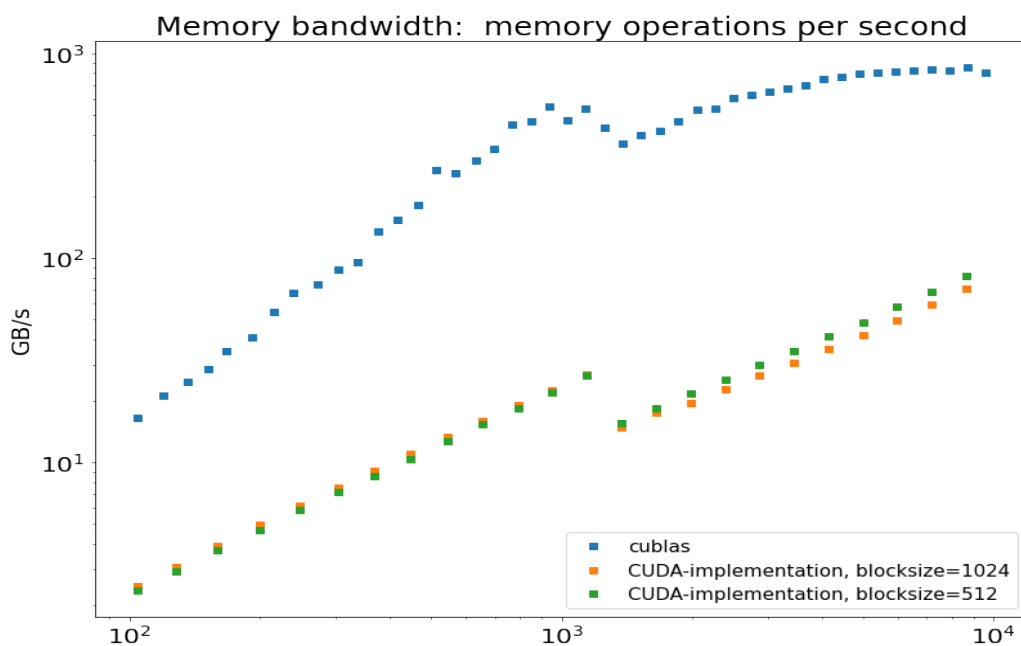
To measure computational performance for matrix-matrix product, we use $2 \times M \times N \times K$ factor.

Simple parallel matrix-vector implementation.

In this task, we implement the simplest matrix-vector function:


```
__global__ void compute_matrix_vector_product(const int M, const int N, const float *x, const float *y, float *z)
{
    const int idx = threadIdx.x+blockIdx.x*blockDim.x;
    if (idx < M)
    {
        z[idx] = 0;
        for (unsigned int j = 0; j < N; j++)
            z[idx] += (x[M*j+idx] + y[j]);
    }
}
```

First, we perform computing for square matrices. We compare our performance with CUBLAS library functions.



The implementation of product has much lower performance in comparison to Cublas. I suppose, to achieve more performance, the loop in kernel function also can be paralleled. The maximum achieved memory bandwidth by Cublas is 852.994 GB/s while my is 81.6256GB/s.

Next, we perform calculation for cases when either number of rows or number of columns is fixed. First, we fix $N = 10000$ and vary M . After, we fix $N = 16384$ and vary M .



For large fixed number of rows, we have quite good performance (because in our implementation we paralleled over rows).

Simple Transponation

Here we implement the naive transponation:

```
__global__ void transposeSimple(const int N, const int M, const float *x, const float *y)
{
    unsigned int xIndex = blockDim.x * blockIdx.x + threadIdx.x;
    unsigned int yIndex = blockDim.y * blockIdx.y + threadIdx.y;
    if (xIndex < N && yIndex < M)
    {
        y[xIndex+N*yIndex] = idata[yIndex+M*xIndex];
    }
}
```

This implementation lost a lot of performance because does not use localization of data in cache memory. But it is really easy to implement.

We compute product $A^T x$ for $M=N=5000$ using this implementation and CUBLAS. CUBLAS implementation showed 783.068 GB/s and mine 42.3618 GB/s.

Matrix-Matrix product implementation.

Here, we Implement matrix-matrix product for case $M=N=K$.

```
__global__ void compute_matrix_matrix_product(const int N, const float *x, const float *y, float *z)
{
    int i=threadIdx.x+blockIdx.x*blockDim.x;
    int j= threadIdx.y+blockIdx.y*blockDim.y;
    float value = 0;
    if (i<N && j<N)
    {
        for ( int k=0;k<N; k++)
            value +=x[i*N+k]*y[k*N+j];

        z[i*N+j] = value;
    }
}
```

We measure computing performance.

