

Targeting the GPU with Kokkos

Kamilla Faizullina

In this work, we implement a computation of the element matrices with Kokkos and run the program on GPU and CPU to check the performance. Each element matrices is computed using its own Jacobian matrices (which are all the same, in our case). Therefore, we assume that each element computation is independent and we can use parallelism along axis linked to Number of elements.

We use NVIDIA Tesla V100 Nvlink. Memory bandwidth is 900 GB/s and the performance is 7.8 teraFLOPS for Double-Precision and 15.7 teraFLOPS for Single-Precision.

1 Implementation

The matrices J (the inverse and transpose Jacobian) and A (results) are stored as Kokkos::View type:

```
Kokkos::View<Number *[3][3], LayoutType, MemSpace> J("J", N);  
Kokkos::View<Number *[4][4], LayoutType, MemSpace> A("A", N);
```

Here 3 dimensions in arrays: run-time, 2 compile-time.

We have $3 \times 3 + 4 \times 4 = 25$ elements for two arrays for every computation. Therefore, I use the factor $25N$ when calculate the performance.

2 Performance for CPU and the GPU

First, we run the program on GPU and CPU (for 4 and 16 processes) with default LayoutType. According to documentation , for execution space Cuda, the default layout is *LayoutLeft*. When the execution space is OpenMP, the layout is *LayoutRight*.

We measure the performance for both float and double types arrays.

Figure 1 illustrates the obtained results. As we see, results from GPU are much better than from CPU on 16 processes. The computational performance for type *float* are better than for *double* in all cases.

For large N , the obtained memory bandwidth on GPU with floats is approximately 650 – 750 GB/s, which is close enough to the theoretical technical value 900 GB/s (also, maybe calculation $25N$ does not take into account some other made actions).




Figure 1: Measured performance for GPU and CPU (for 4 and 16 processes) using Kokkos with default layouts.

3 LayoutRight vs LayoutLeft

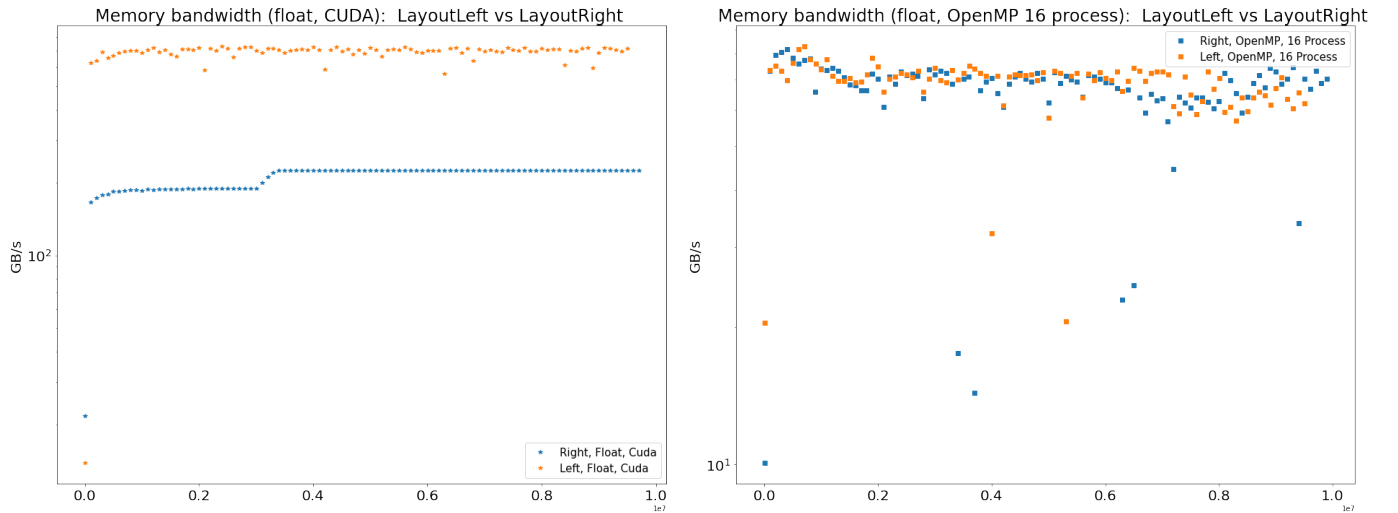


Figure 2: Measured performance for GPU and CPU (for 16 processes) using Kokkos with LayoutLeft and LayoutRight.

Kokkos allows to choose manually the Layout in Arrays. LayoutLeft is column-major layout and LayoutRight is row-ordered. Figure 2 shows how choice of Layout affect the performance. As we can see, for GPU the leftlayout choice is better. Therefore, it is not surprising that for simple parallels the LayoutLeft is default choice on Cuda. For runs with OpenMP, the difference is not significant. However, it was expected that the LayoutRight would outperform.

4 Timing

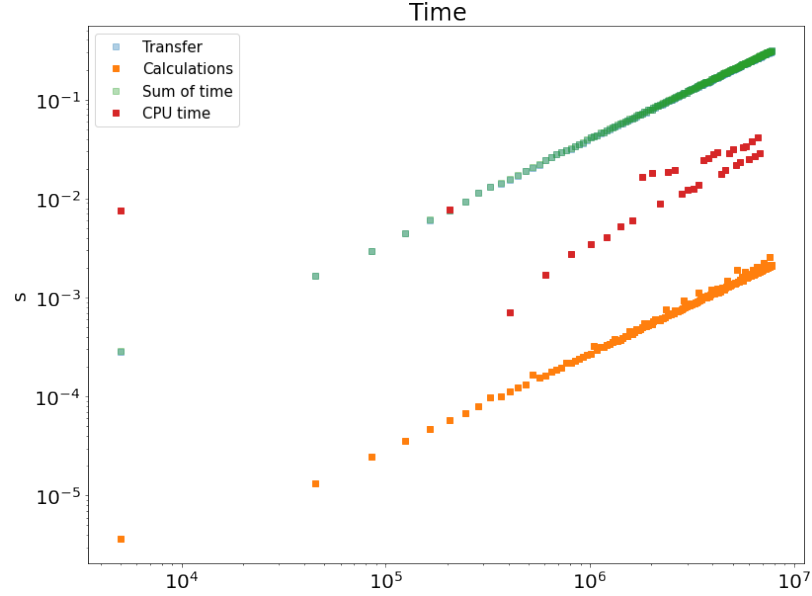


Figure 3: Measured time for transfer and calculations.

Figure 3 shows how much time transferring data between host and device costs for GPU. The total time is mostly time for copy. The difference approximately 100 times. In general, running CPU was less time-consuming because there is no data transfer.

I suppose that for one-time operations we should use CPU to save time on transferring. For iterative actions, we should use GPU (however, it depends on the number of iterations). if iterations less than 100 I am not sure that we would enough.